


1) Intro -- Deep Generative Models

- **Generative Model:** Model a generative process that defines a joint distribution over random variables and their stochastic interactions.
- **Usage:** *Data Generation (sampling)*, *Density estimation (anomaly detection)*, *Unsupervised representation learning (Learn common features)*
- **Product rule:** Allows us to *factorize* the joint distribution in two manners -- $p(x, y) = p(x|y) p(y) = p(y|x) p(x)$
(Joint distribution could be represented as a product of a marginal distribution and a conditional distribution)
- **Sum rule:** If want calculate marginal distribution over one of the variables, must integrate/sum out other variable -- $p(x) = \sum_y p(x, y)$

- **Deep G.M vs G.M:** Distribution is parameterized using deep neural networks (Rich and flexible parameterization of distributions)

- Continuous Random Variable:

- If X is a continuous random variable, we can usually represent it using its probability density function (PDF) p_x
- If X is a continuous **random vector**, we can usually represent it using its joint probability density function:




- Discriminative approach VS Generative approach:

- *Discriminative approach model only $p(y|x)$... Called a discriminative model because it is only useful for discriminating Y 's label when given X*
- *Problem: Adding noise to images could result in completely false classification...*
- *NNs used to parameterize conditional distribution $p(y|x)$ seem to lack understanding, and no uncertainty quantification...*
- *Generative approaches are key for models that understand and quantify beliefs (and for that, **estimating $p(x)$** is crucial!!)*

- Explicit vs Implicit Generative Models:

- **Explicit:** Define a density function of the distribution
- **Implicit:** Learn mapping that generates samples by transforming an easy-to-sample random variable




$$p(x, y) = p(x|y) p(y) \text{ VS } p(x, y) = p(y|x) p(x)$$

- Now, the following models answers the key question of **how to represent $p(x)$** :

2) Autoregressive Models

- Autoregressive approach casts $p(x)$ as a product of conditional distributions...
- Define the joint distribution using conditionals over each feature, given the values of the previous features.
- The distribution over x is represented in an autoregressive manner: $p_\theta(x) = \prod_{i=1}^d p_\theta(x_i | x_{<i})$, where $x_{<i}$ denotes all x 's up to the index i .
- Main idea: Leverage the chain rule of probability to factorize the joint probability distribution over the input space into a product of conditional distributions
- Each conditional probability is parameterized by a deep neural network whose architecture is chosen according to the required inductive biases for the data.
- Big challenge is to calculate these conditional likelihoods $p(x_i | x_1, \dots, x_{i-1})$.
- How can we define these complex distributions in an expressive model that is also tractable and scalable? One solution is to use universal approximators (ex: DNN)
- Drawback: Parameterized in an autoregressive manner, hence, sampling is rather a slow process
- Example: PixelCNN
 - A pixel of interest i (in red) is defined by all previous pixels (in blue).
 - The PixelCNN can model their associations using convolutional layers along a deep neural network.




3) Normalizing Flows

- Technique for building complex probability distributions by transforming simple ones
- Transform the simple distribution with some function f
- f represents a composition of a sequence of invertible transformations that we are trying to learn

- We are mapping x to its inverse z
- Evaluating the density of this z under $p(z)$
- Then multiplying by some scalar magnitude (Space Expansion/Contraction indice)
- The change of variables formula provides a principled manner of expressing a density of a random variable by transforming it with invertible transformation f .
- $p_\theta(x) = p(z) |J_{f(z)}|$, where $z = f_\theta(x)$ and $J_{f(z)}$ denotes the Jacobian matrix.

- Learn a Sequence of invertible transformations ($f_k: \mathbb{R}^D \rightarrow \mathbb{R}^D$)
- Start with a simple known distribution
- Then sequentially apply invertible transformations to obtain flexible distribution

- Each f is invertible (bijective)
- $x = f_\theta(z) = f_k \circ \dots \circ f_2 \circ f_1(z)$ $f: Z \rightarrow X$ and $J_{f(z)}$




- Learning objective: log-likelihood function (Maximum likelihood objective where we directly maximize the density assigned to x by the model)
- Sampling via forward transformation ($z \rightarrow x$): $z \sim p_z(z)$ $x = f_\theta(z)$
- Latent representations inferred via inverse transformation: $z = f^{-1}_\theta(x)$

- Role of magnitude of the determinant of the Jacobian of F inverse:
 - Indicates how much transformation locally expands/contracts space
 - Necessary to ensure new density function $p(x)$ satisfy requirement $\int p(x) dx = 1$
 - Hence the name "normalized" flows...

4) Latent Variable Models

A) Variational Autoencoder (VAE)

- **Encoder:** Produces a distribution over latent variables z given a data input x
- **Decoder:** Reconstructs the data by producing a distribution over data \hat{x} given a latent input z



- **Training Objectives:**
 - x observed variables z latent variables
 - Variational Lower Bound (Evidence Lower bound)
 - Lower bound on the marginal log likelihood $p_\theta(x)$
 - $\log p_\theta(x) \geq \text{Variational Lower Bound}$

$$\log p_\theta(x) \geq \mathbb{E}_{q(z|x)}[\log p_\theta(x|z)] - D_{KL}(q(z|x) \parallel p_\theta(z))$$


- 1) Likelihood term encourages model to max expected density assigned to the data
- 2) KL divergence term encourages the approximate posterior $q(z|x)$ to be similar to the prior on the latent variable $p(z)$

- (For more details, see book or saved papers)

- Use an explicit density estimation but define an intractable density function with latent variables that cannot be optimized directly.
- To train the model, optimize the lower bound of likelihood instead (approximate density);
- Optimize the log-likelihood of the data by maximizing the evidence lower bound (ELBO)
- All distributions must be defined upfront and, therefore, they are called *prescribed models*.

B) Deep Diffusion Models

- Main idea: Iteratively destroy the structure in data through a forward diffusion process
- Afterward, learn a reverse diffusion process to restore the structure in data.
- *Inspiration from non-equilibrium statistical physics*



- **Two step process:** forward diffusion process and the reverse process (reconstruction)
- 1) Forward diffusion process: gaussian noise is introduced successively until the data becomes all noise. (Markov Chain) (No learning)
- 2) Reverse/ reconstruction process: Undoes noise by **learning the conditional probability densities using a DNN**.



- Reverse process requires estimation of probability density at an earlier time step given the current state of the system.
- Hence generating data sample from isotropic Gaussian noise

- The forward noise step is restricted to be small (low ambiguity about x_{t-1})
- Hence we can model the posterior of the forward step that is $q(x_{t-1}|x_t)$ with a unimodal gaussian (eliminating contributions other modes)
- Shown theoretically that in the limit of infinitesimal step sizes, the true reverse process will have the same functional form as the forward process, which is a key principle used for diffusion models [1949 Feller] where we parameterize each learned reverse step to also be a unimodal diagonal gaussian.

- Unlike the forward process, the estimation of previous state from the current state requires the knowledge of all the previous gradients
- Which we can't obtain without having a learning model that can predict such estimates.
- Therefore, we train a NN that **estimates the $p_\theta(x_{t-1}|x_t)$** based on learned weights θ and current state at time t . (Input: Sample x and I)
- Take t as input also as different timesteps are associated with different noise levels (model can learn to undo these individually)


- Model architecture often simpler compared to other architectures (*Single network vs VAE's 2 networks*)
 - The input layer has the same input size as that of the data dimensions.
 - There can be multiple hidden layers depending on the depth of the network requirements.
 - The middle layers are linear layers with respective activation functions.
 - The final layer is again of the same size as that of the original input layer, thereby reconstructing the original data.

- Training:

- Not maximum likelihood objective (max density assigned to x_0 by model) as if we try to calculate $p(x_0)$, we see that we have to marginalize over all the possible trajectories/ways we could have arrived at x_0 when starting from a noise sample
- As this is intractable, we maximize a lower bound, like the one seen in VAEs

- Latent Variables: (x_1, x_2, \dots, x_T) and Observed Variable: (x_0)

- Forward process in diffusion models as analogous to the VAE encoder producing latency from data (fixed for this case)
- Reverse process as analogous to the VAE decoder producing data from latents
- See links for details about training targets (<https://youtu.be/jbLgFrTrnGU> + books)



- **Sample conditionally given some variable of interest** (ex: class label):
 - One way: Feed the conditioning variable y as additional input during training
 - Model should learn to use y as a helpful hint about what it should be reconstructing in practice

- *Rest coming at the end of this semester (And further details on each models, applications, comparison, ...)*